INF4820: Algorithms for Artificial Intelligence and Natural Language Processing

Introduction and Overview

Stephan Oepen & Murhaf Fares

Language Technology Group (LTG)

August 25, 2016

Topics for Today

Overview

- Course motivation and introduction:
- AI, NLP, ML — What are they?
- Common Lisp — What and why?
- Outline of lectures and learning goals.
- Practical details.

What is AI?

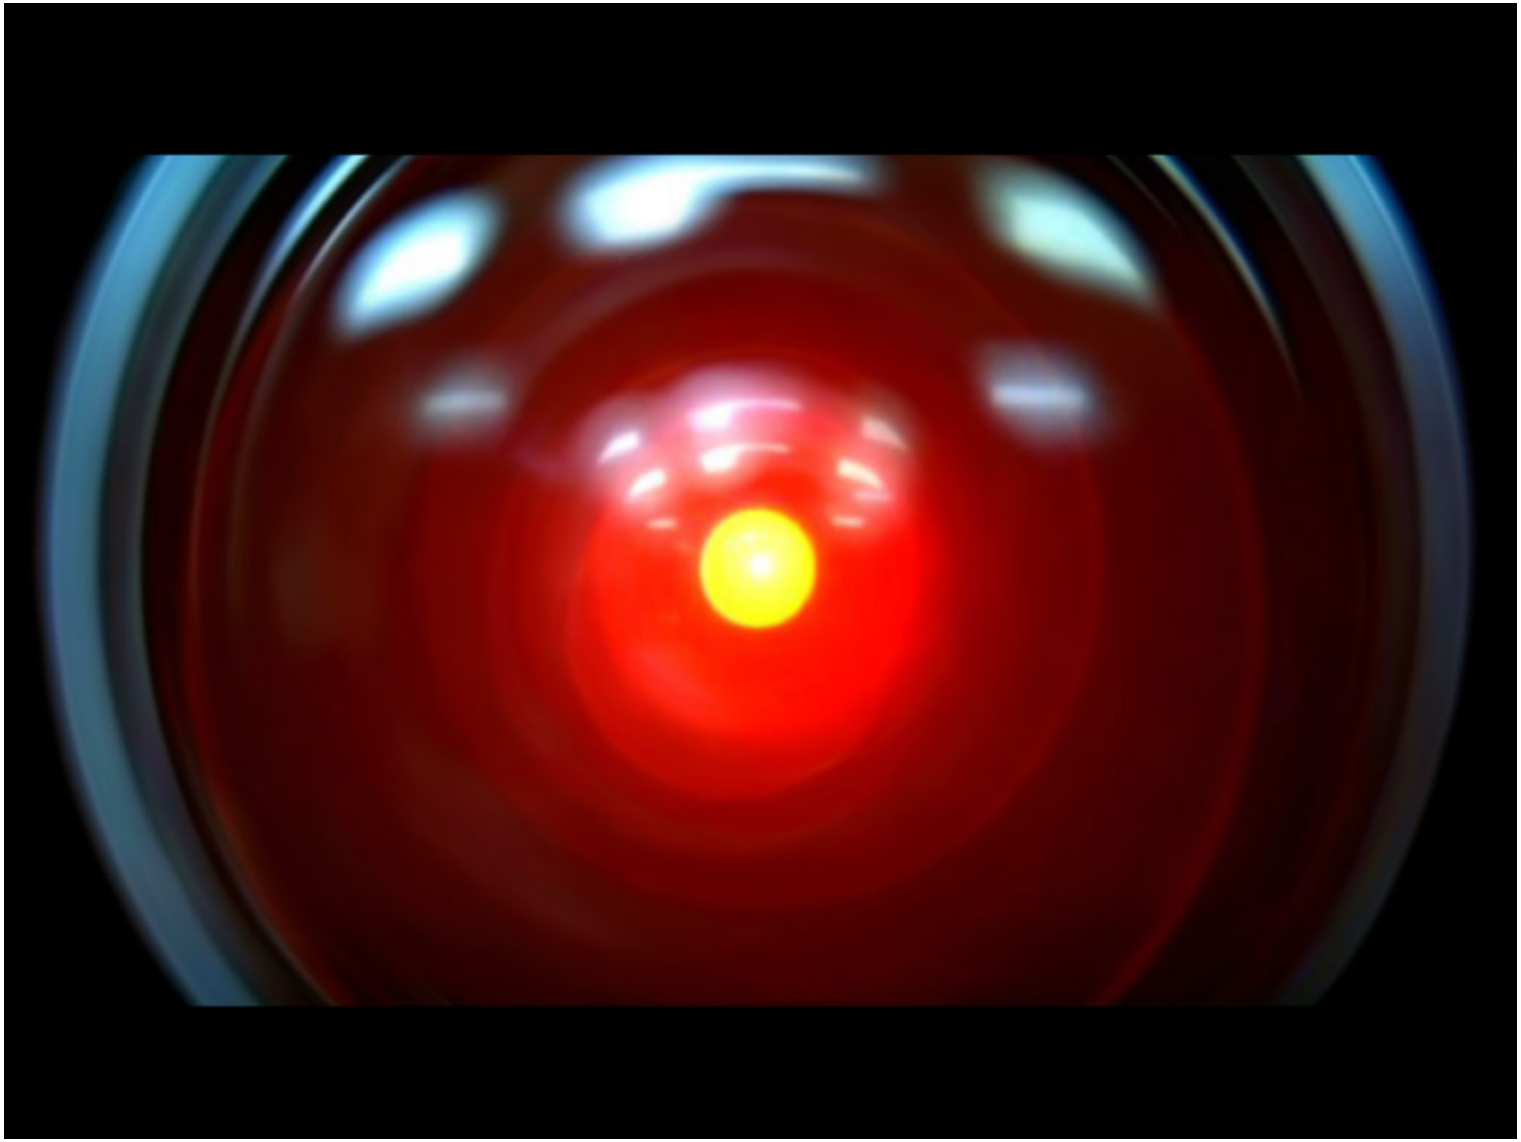

(HAL 9000 in *2001: A Space Odyssey*; 1968)

AI, Hype Cycles, and the End of the World

- AI has always been subject to cycles of hype.
- A lot of media attention recently concerning the existential risk of AI:
- Stephen Hawking et al. in an op-ed in the Independent (spring 2014):
  - *it's tempting to dismiss the notion of highly intelligent machines as mere science fiction. But this would be a mistake, and potentially our worst mistake in history.*
  - *Whereas the short-term impact of AI depends on who controls it, the long-term impact depends on whether it can be controlled at all.*
- Elon Musk (Tesla Motors, SpaceX) at a MIT talk (fall 2014):
  - *With AI we are summoning the demon.*

The Future of Life Institute

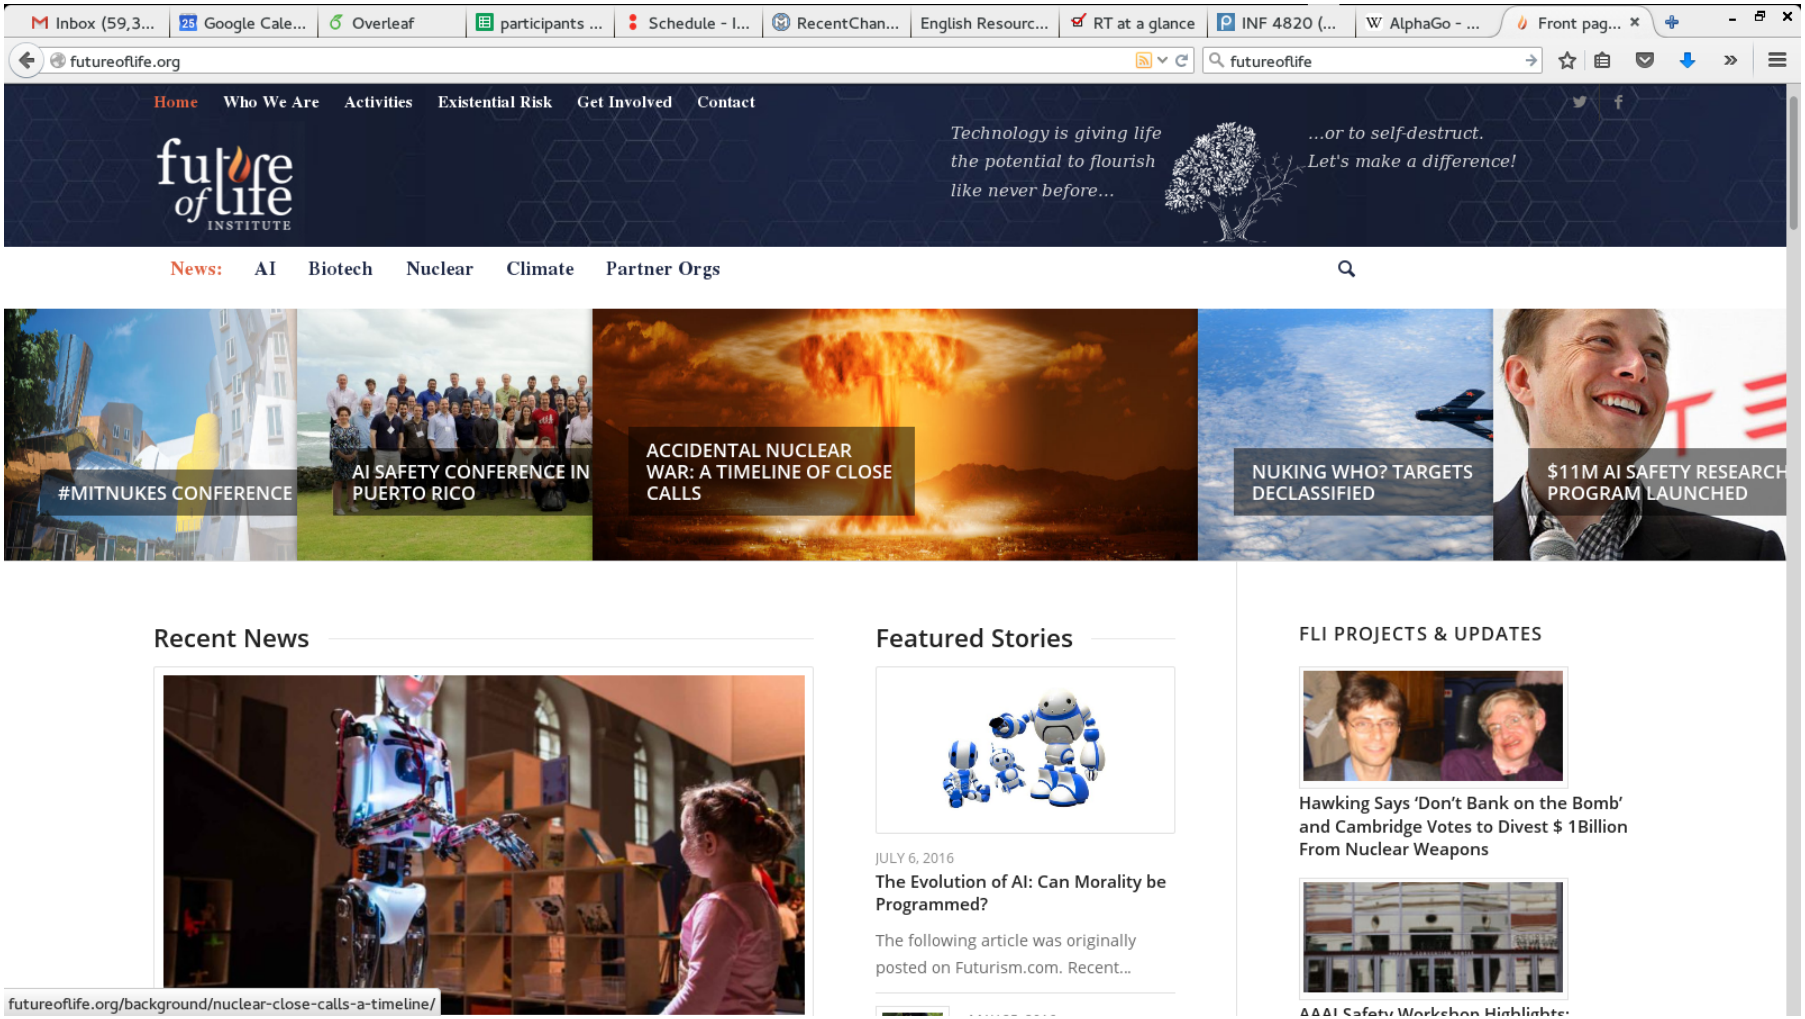

(Spring 2015: Alarm by Hawking, Musk, Norvig, Russel, Wozniak, etc.)

AI in Our Daily Lives (1 of 3)

Search r... | Google ... | Overleaf | participa... | Managing: in... | RecentC... | English Res... | RT at a g... | Managing: in... | Piazza | ... | self-drivi... | Tesla ... | The first ...

https://www.theguardian.com/technology/2016/jun/30/tesla-autopilot-death-self-driving-car-elon-musk

self-driving car death

Tesla

Tesla driver dies in first fatal crash while using autopilot mode

The autopilot sensors on the Model S failed to distinguish a white tractor-trailer crossing the highway against a bright sky

Danny Yadron and Dan Tynan in San Francisco

Friday 1 July 2016 00.14 BST

This article is 1 month old

Shares
9,480

Save for later

Joshua Brown, the first person to die in a self-driving car accident. Photograph: Facebook

The first known death caused by a self-driving car was disclosed by Tesla Motors on Thursday, a development that is sure to cause consumers to second-guess the trust they put in the booming autonomous vehicle industry.

The 7 May accident occurred in Williston, Florida, after the driver, Joshua Brown, 40, of Ohio put his Model S into Tesla's autopilot mode, which is able to control the car during highway driving.

Advertisement Report this ad

Thanks for the feedback! Undo

We'll use your feedback to review ads on this site.

(July 2016: A Tesla driver dies while auto-pilot was in control)

AI in Our Daily Lives (2 of 3)

(July 2016: Minor injuries from run-in with security robot)

AI in Our Daily Lives (3 of 3)

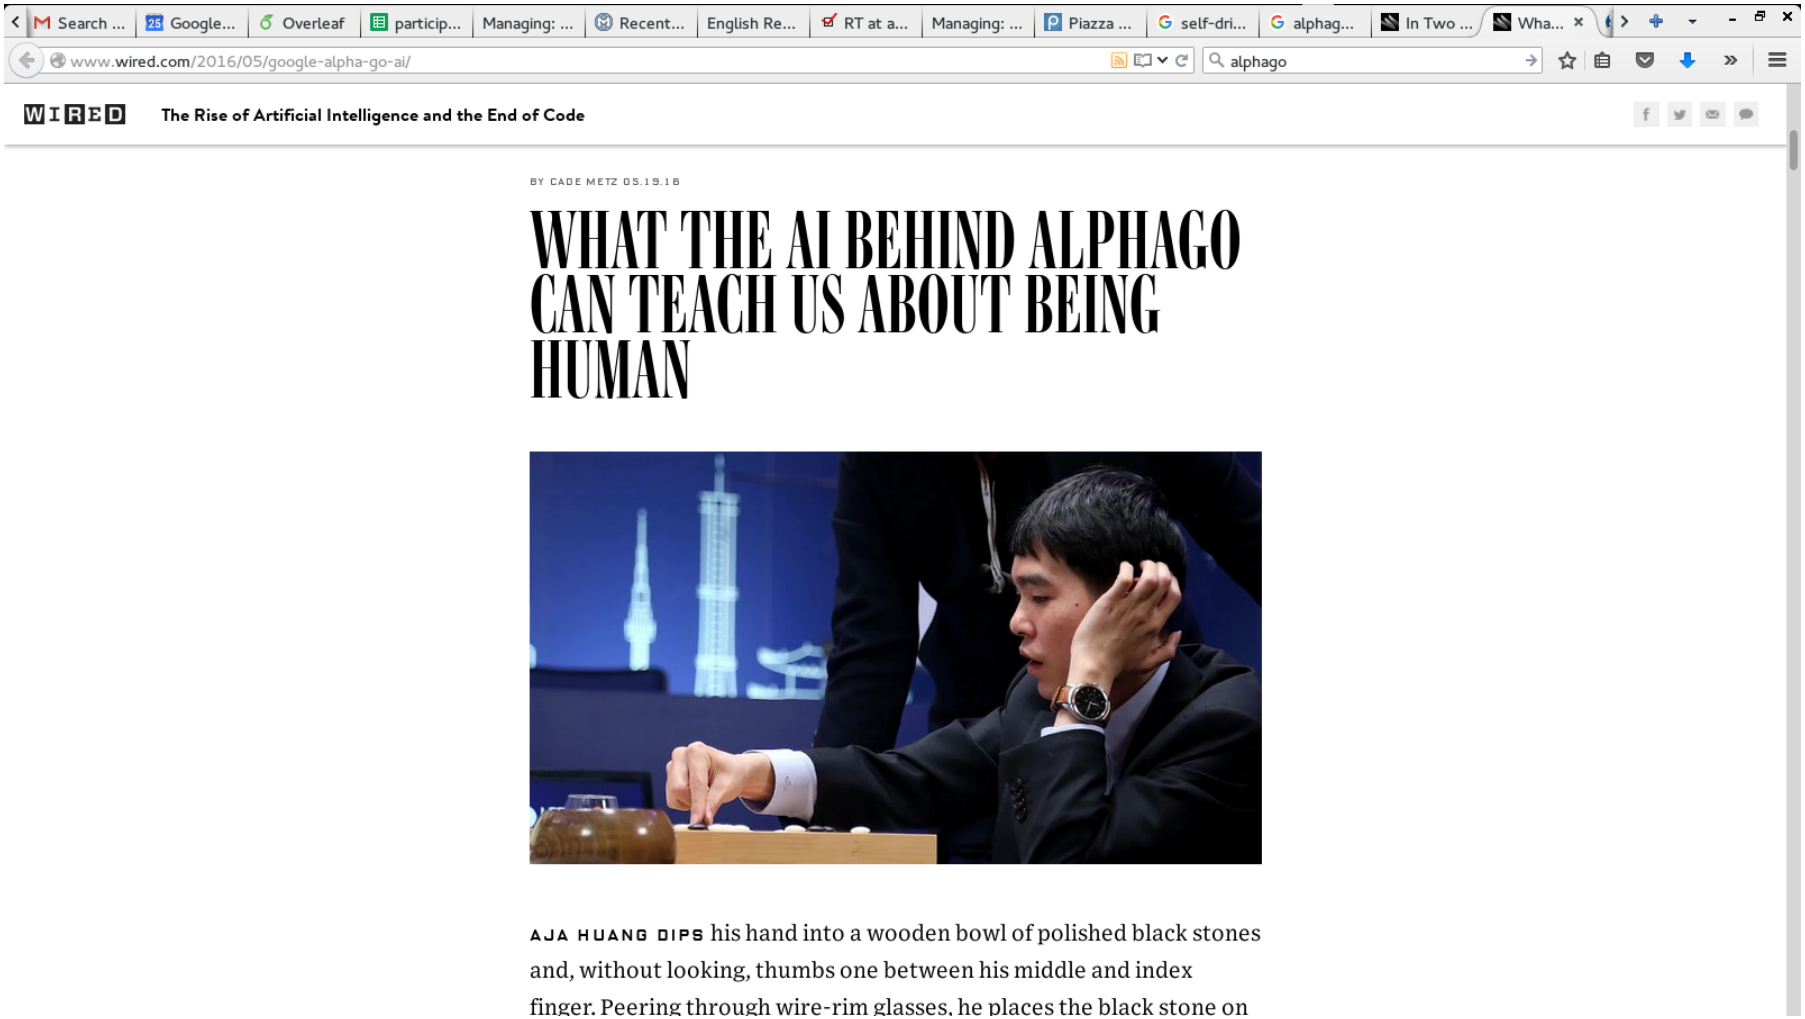

(March 2016: Nine-dan professional Go player Lee Sedol looses 4–1)

What is AI?

- Alan Turing, 1950:
  - *I propose to consider the question, 'Can machines think?'*
- The term 'AI' coined by John McCarthy (Dartmouth Conference, 1956).
  - *The science and engineering of making intelligent machines.*
  - *Every aspect of learning or any other feature of intelligence can be so precisely described that a machine can be made to simulate it.*
- Language always in central place, cf. the Turing Test.

What is AI? (cont'd)

- The early years: simple chatbots, theorem proving, blocks world, expert systems, game playing (chess), ...
- Moving target: Whatever requires 'intelligent' decisions, but seems out of reach technologically?
  - Web search arguably would have been AI a couple of decades ago.
  - Open-domain Machine Translation out of reach until around 2005.
- For our purposes: AI is a toolkit of methods for representation and problem solving.

What is Natural Language Processing?

- Making computers 'understand' human language
- Aka language technology or computational linguistics
- Young and interdisciplinary field:
- Computer Science + Linguistics
- (+ Cognitive Science + Statistics + Information Theory + Machine Learning + . . .)

Some Applications

- Grammar and/or spell checkers, auto-completion
- Machine translation
- Q&A systems
- Dialog systems
- Speech recognition and synthesis
- Intelligent information extraction
- Summarization
- Sentiment analysis
- Any application requiring an understanding of language. . .

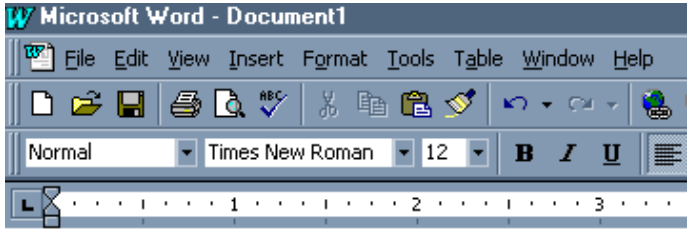

What is AI?

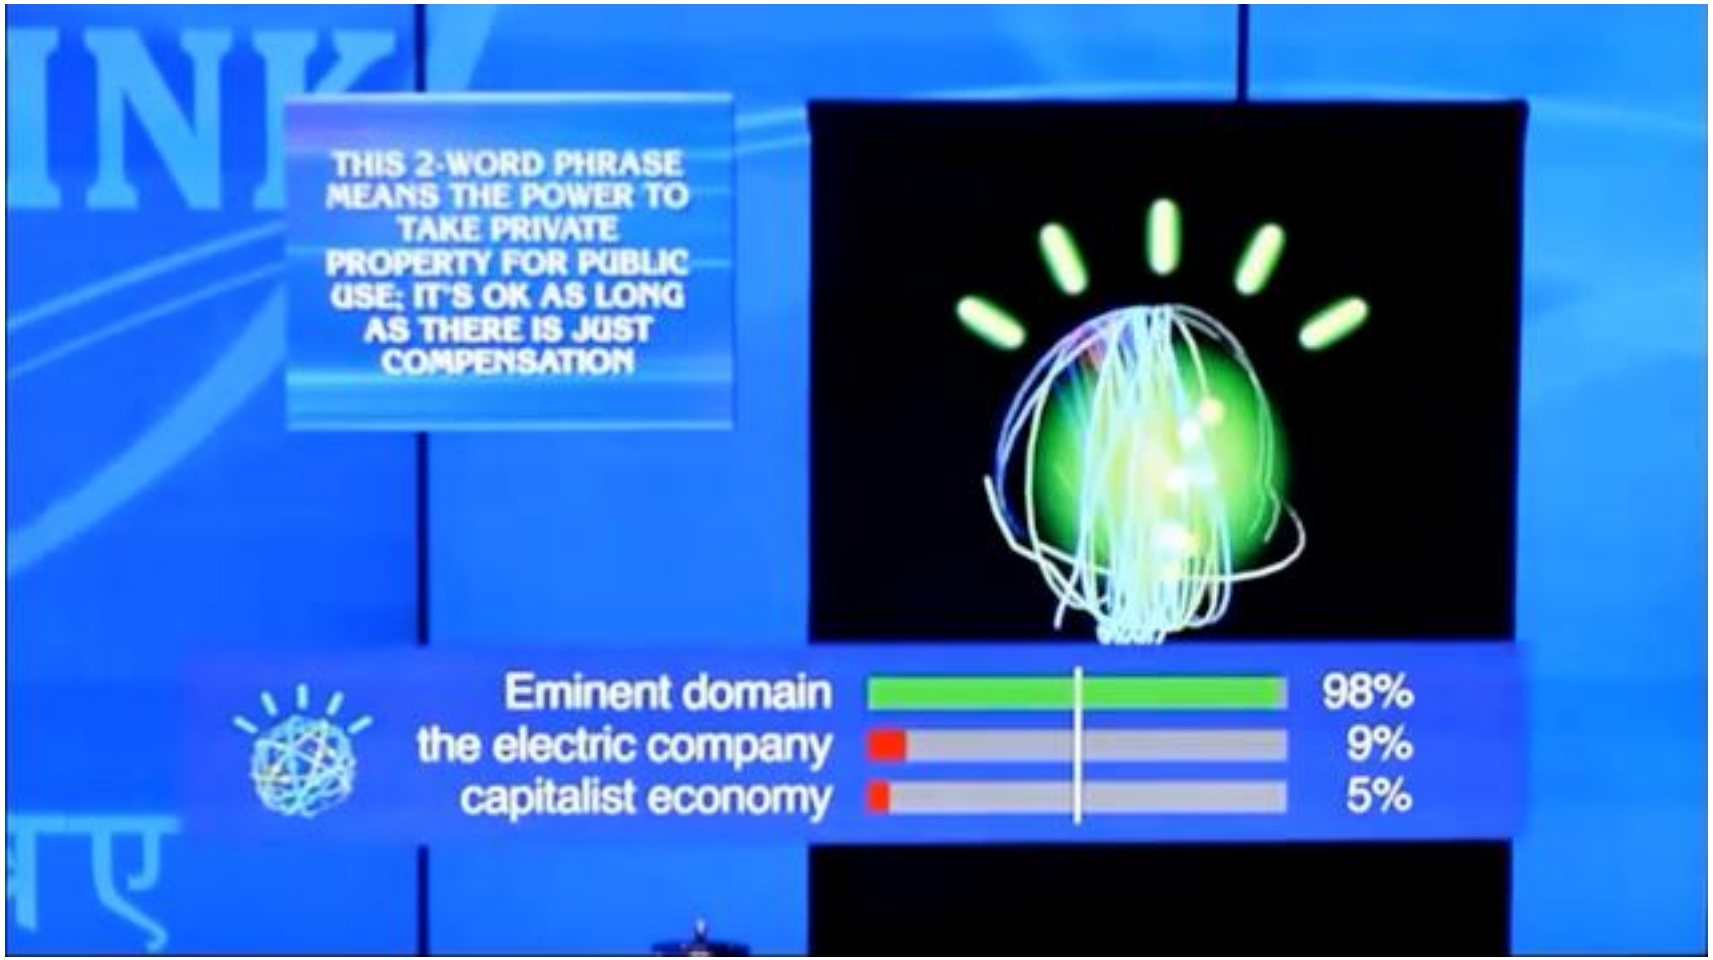

(IBM Watson beats long-time *Jeopardy!* champions; 2011)

What Makes NLP Hard?

Ambiguity

- I.e. the property of being open to multiple interpretations.
- All levels of linguistic description are associated with ambiguities.
- For humans, ambiguity is a feature: language is an efficient code.
  - The same expressions can be re-used in different contexts.
  - A large part of the information can be underspecified.
  - Interpretation relies on background knowledge and our expectations in a given context of use.
- Disambiguation is a central problem in NLP $\longrightarrow$ Search problems.

Ambiguity: Some examples

Word level ambiguity

- Norwegian: *rett*.
- English: *meal*, *dish*, *straight*, *correct*, *fair*, *justice*, *right*, *court*, *law*, *direct*, *grade*, . . . ?
- Ambiguous in meaning + syntactic category (part of speech).
- Need context to decide.

De hadde laget en deilig rett av grønnsaker.

Streken må være rett.

Kunden har alltid rett.

Du har rett til en advoktat.

Det er lovlig i henhold til norsk rett.

Slikt skjer rett som det er.

Vennligst rett disse prøvene!

Vi kjørte rett hjem.

Ambiguity: Some Examples

Referential Ambiguity

The authorities jailed the protesters because they { *advocated revolution.* / *feared revolution.*

Sentence-Level Ambiguity

I like eating sushi with { *tuna.* / *sticks.*

Acoustic Ambiguity

Let's talk about how to { *recognize speech* / *wreck a nice beach*

Some History: Different Approaches to AI and NLP

- Traditionally; two broad paradigms in NLP (and AI).
  - The rationalist approach, based on hand-crafted formal rules and manually encoded knowledge.
  - The empiricist approach, based on automatically inferring statistical patterns from data.
- 1950s – 80s: Rule-based
- Late 1980s: Empirical systems outperform rule-based in the area of speech recognition.
- 1990s: NLP as whole sees a shift of interest from rationalist towards empirical approaches.
- 2000s: No longer conceived as opposing poles, but complementary approaches typically used together.

The Basics of Empirical Methods

- The theoretical foundations are studied within the field of machine learning (ML) or statistical learning theory.

Machine Learning

... *the study of computer algorithms that improve automatically through experience* (Tom Mitchell 1997).

- Goal: Learn from examples, to make predictions about new data.
- Has applications in many other data-intensive sciences besides NLP, e.g. meteorology, biology, physics, robotics, signal processing, etc.
- An arsenal of methods: decision trees, support vector machines, maximum entropy models, naïve Bayes classifiers, artificial neural networks, genetic algorithms, ...

Lisp

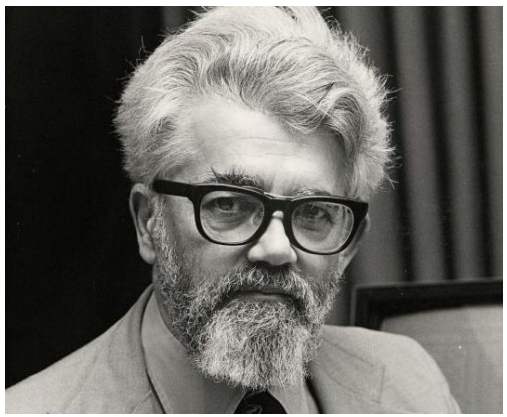

- Powerful high-level language with long traditions.
- Especially strong support for symbolic and functional programming.
- "Discovered" by John McCarthy in 1958.
  - Initially intended as a mathematical formalism.
  - Then one of his students, Steve Russell, implemented an interpreter for the formalism, and Lisp the programming language was born.
- Rather than Lisp becoming outdated, the tendency has been that other languages have developed towards Lisp.

Common Lisp

```
(print "Hello world!")
```

- Several dialects; we will be using Common Lisp.
- Fully ANSI-standardized and stable.
- Rich language: multitude of built-in data types and operations.
- Easy to learn:
  - extremely simple syntax;
  - straightforward semantics.

An Experiment in Live Programming

The Factorial Function

$$n! \equiv \begin{cases} 1 & \text{for } n = 0 \\ n \times (n-1)! & \text{for } n > 0 \end{cases}$$

Common Lisp Implementation

```
(defun ! (n)
  (if (= n 0)
    1
    (* n (! (- n 1)))))
```

A Note on Lisp and AI

- Often hailed (or dismissed) as "the AI language".
- While not quite true, there are several reasons for this coupling:
- AI coined by McCarthy in the mid-1950s.
- Lisp conceived by McCarthy in the mid-1950s.
- In addition to being fast and powerful, Lisp is particularly well suited for:
  - Explorative programming
  - Rapid prototyping
  - Incremental and interactive development
  - Extending the language itself

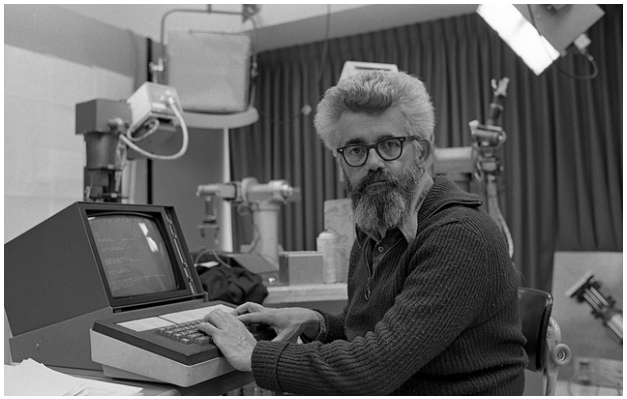

Lisp + Emacs = Good Match

- Steep learning curve, but with a big pay-off:
- Emacs is an unusually powerful editor.
- Written in Emacs Lisp.
- Highly customizable—the Emacs Lisp dialect is also used as an extension language.
- Different “modes” make Emacs sensitive to different editing needs, e.g. depending on the specific programming language used.
- Prerequisite for an enjoyable Emacs experience: Spend some time mastering basic key commands!

Overview of Lectures

- Common Lisp basics
- Vector space models
- Classification and clustering
- Probability theory
- Hidden Markov Models
- Statistical parsing

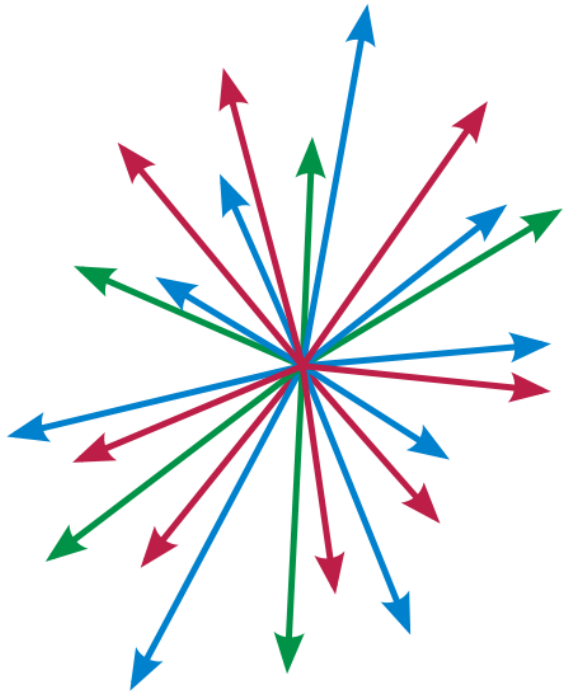

- Recurring themes: Machine learning, scalable data representations, search, dynamic programming.
- Two hours of lectures every week; two-hour laboratory weekly

Very High-Level Course Summary

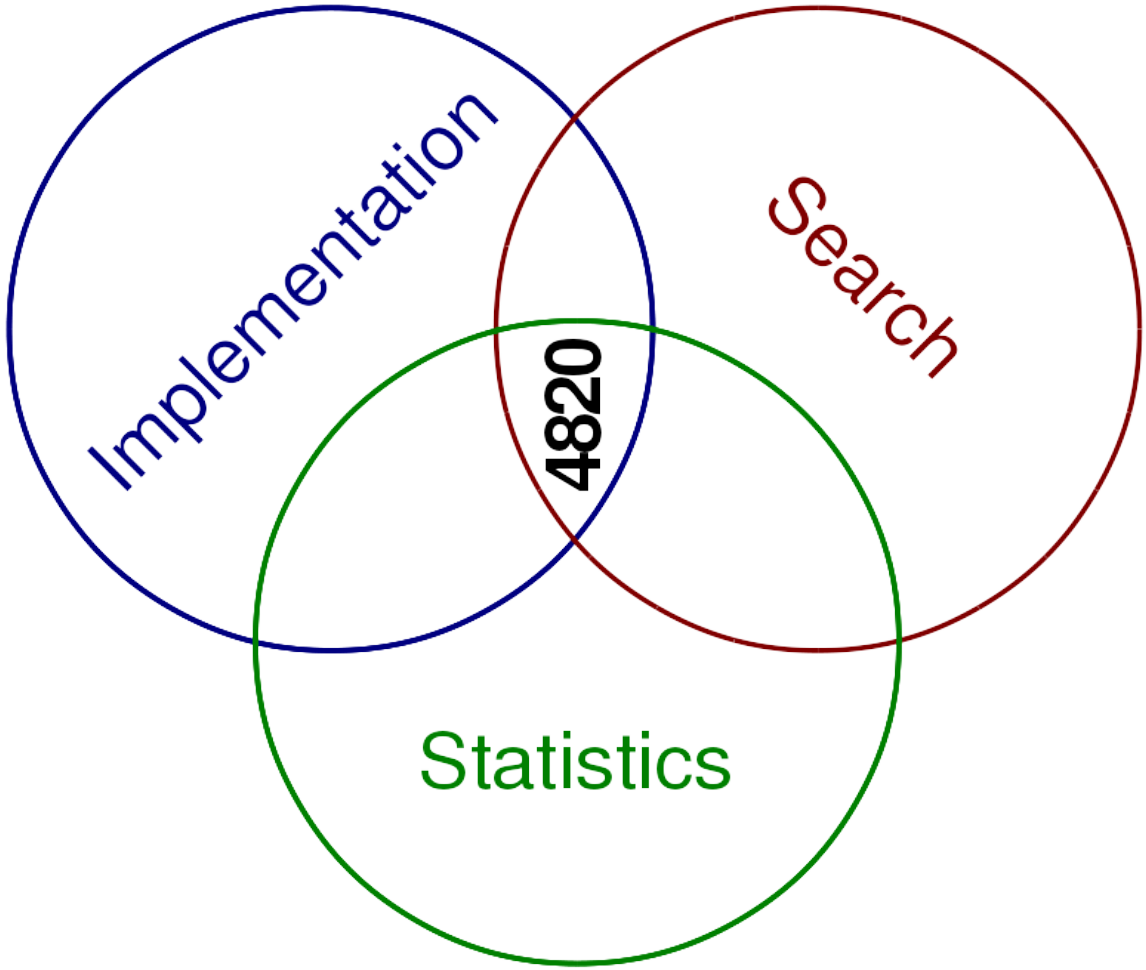

Efficient and Scalable Algorithms and Data Structures for Searching (Probabilistically) Weighted Solution Spaces

Obligatory Exercises

- Three obligatory exercises:
- Exercise (2) and (3) have two parts each;
- Five problem sets in total.
- In order to pass and qualify for the exam you need a least
  - 6 of 10 possible points for Exercise (1),
  - 12 of 20 possible points for (2a) + (2b),
  - 12 of 20 possible points for (3a) + (3b).
- Extensions can only be given in case of illness, and re-submissions will not be possible.
- See course page for the schedule (tba).

Infrequent Experiments in High-Tech Teaching

- For student involvement and incremental exam preparation:
- occassional short quiz sessions → extra points towards exercises.

Example Quiz (0 + 0 Points)

1. Live programming can be useful?

   A: yes; B: no

2. Lisp was first developed by:

   A: Alan Turing; B: John McCarthy

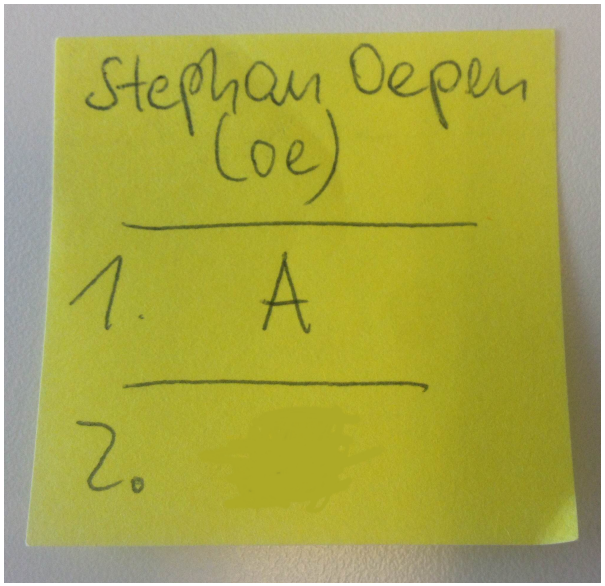

Reading List

Obligatory reading; *selected parts* from:

- Jurafsky & Martin (2008):
  Speech and Language Processing (2nd Ed.)
- Seibel (2005):
  Practical Common Lisp (Available On-Line)
- Manning, Raghavan, & Schütze (2008):
  Introduction to Information Retrieval (Available On-Line)

Other recommended resources:

- Despite being 20 years old and long out-of-print *On Lisp* by Paul Graham is still a great read.
  - Freely available on-line: `http://www.paulgraham.com/onlisp.html`
- The Common Lisp 'HyperSpec':
  - `http://www.lispworks.com/documentation/HyperSpec/Front/`

Contact

- Questions?
  - Piazza: on-line discussion board linked from course page.
  - `inf4820-help@ifi.uio.no` reaches all course staff:
  - Murhaf Fares;
  - Stephan Oepen;
  - Elena Volkova (laboratory assistant);
  - `{murhaff|oe|elenavo}@ifi.uio.no`.
- Messages:
  - Check your UiO email regularly;
  - Subscribe to the RSS feed of the course page;
  - Participate in the on-line discussion board.